

Rats - An Bayesian analysis

Ivan Ferrante (1563390)

```
library(R2jags)
library(ggplot2)
```

Introduction

The data contains 30 young rats whose weights were measured weekly for five weeks. Dependent variable Y_{ij} is the weight of the i^{th} rat at age x_j .

```
rats.data <- list(x = c(8.0, 15.0, 22.0, 29.0, 36.0), xbar = 22, N = 30, T = 5,
  Y = c(151, 199, 246, 283, 320,
        145, 199, 249, 293, 354,
        147, 214, 263, 312, 328,
        155, 200, 237, 272, 297,
        135, 188, 230, 280, 323,
        159, 210, 252, 298, 331,
        141, 189, 231, 275, 305,
        159, 201, 248, 297, 338,
        177, 236, 285, 350, 376,
        134, 182, 220, 260, 296,
        160, 208, 261, 313, 352,
        143, 188, 220, 273, 314,
        154, 200, 244, 289, 325,
        171, 221, 270, 326, 358,
        163, 216, 242, 281, 312,
        160, 207, 248, 288, 324,
        142, 187, 234, 280, 316,
        156, 203, 243, 283, 317,
        157, 212, 259, 307, 336,
        152, 203, 246, 286, 321,
        154, 205, 253, 298, 334,
        139, 190, 225, 267, 302,
        146, 191, 229, 272, 302,
        157, 211, 250, 285, 323,
        132, 185, 237, 286, 331,
        160, 207, 257, 303, 345,
        169, 216, 261, 295, 333,
        157, 205, 248, 289, 316,
        137, 180, 219, 258, 291,
        153, 200, 244, 286, 324))
```

```
Y <- matrix(rats.data$Y, nrow = rats.data$N, ncol = rats.data$T, byrow = TRUE)
Y
```

```
##      [,1] [,2] [,3] [,4] [,5]
## [1,] 151 199 246 283 320
## [2,] 145 199 249 293 354
## [3,] 147 214 263 312 328
## [4,] 155 200 237 272 297
```

```
## [5,] 135 188 230 280 323
## [6,] 159 210 252 298 331
## [7,] 141 189 231 275 305
## [8,] 159 201 248 297 338
## [9,] 177 236 285 350 376
## [10,] 134 182 220 260 296
## [11,] 160 208 261 313 352
## [12,] 143 188 220 273 314
## [13,] 154 200 244 289 325
## [14,] 171 221 270 326 358
## [15,] 163 216 242 281 312
## [16,] 160 207 248 288 324
## [17,] 142 187 234 280 316
## [18,] 156 203 243 283 317
## [19,] 157 212 259 307 336
## [20,] 152 203 246 286 321
## [21,] 154 205 253 298 334
## [22,] 139 190 225 267 302
## [23,] 146 191 229 272 302
## [24,] 157 211 250 285 323
## [25,] 132 185 237 286 331
## [26,] 160 207 257 303 345
## [27,] 169 216 261 295 333
## [28,] 157 205 248 289 316
## [29,] 137 180 219 258 291
## [30,] 153 200 244 286 324
```

```
T <- rats.data$T
T
```

```
## [1] 5
```

```
x <- rats.data$x
x
```

```
## [1] 8 15 22 29 36
```

```
xbar <- rats.data$xbar
xbar
```


```
## [1] 22
```

```
N <- rats.data$N
N
```

```
## [1] 30
```

We can plot each rats growth in the time, by pick the transpose matrix of our data:

```
times = as.numeric(rats.data$x)
matplot(times, t(Y), type = "b", pch = 16, lty = 1)
```



We can use the frequentistic approach in order to predict the weight Y from time x_j for each rat.

```
# LINEAR REGRESSION
array_intercept <- rep(NA,N)
array_slope <- rep(NA,N)
for(i in 1:N)
{
  linear_reg <- lm(Y[i,] ~ x)
  array_intercept[i] <- linear_reg$coefficients[[1]]
  array_slope[i] <- linear_reg$coefficients[[2]]
}
array_intercept

## [1] 107.17143 87.08571 108.22857 120.31429 84.11429 114.22857 98.08571
## [8] 105.91429 123.88571 92.05714 105.11429 93.40000 106.94286 118.65714
## [15] 128.71429 116.85714 93.20000 114.05714 111.82857 109.28571 106.42857
## [22] 97.94286 104.48571 117.60000 77.37143 107.94286 126.88571 116.65714
## [29] 95.68571 106.88571

array_slope

## [1] 6.028571 7.314286 6.571429 5.085714 6.685714 6.171429 5.914286
## [8] 6.485714 7.314286 5.742857 6.985714 6.100000 6.157143 6.842857
## [15] 5.185714 5.842857 6.300000 5.742857 6.471429 6.014286 6.471429
## [22] 5.757143 5.614286 5.800000 7.128571 6.657143 5.814286 5.742857
## [29] 5.514286 6.114286

# INTERCEPT MEAN
mean_alpha <- mean(array_intercept)
mean_alpha


## [1] 106.5676

# SLOPE MEAN
mean_beta <- mean(array_slope)
```

```
mean_beta
```

```
## [1] 6.185714
```

```
plot(x,colMeans(Y), lwd=4, xlab = "days", ylab = "weight",  
     col="red", ylim=c(135,355))  
points(rep(x[1],N), Y[,1])  
points(rep(x[2],N), Y[,2])  
points(rep(x[3],N), Y[,3])  
points(rep(x[4],N), Y[,4])  
points(rep(x[5],N), Y[,5])  
abline(mean_alpha, mean_beta, col="orchid", lwd=2.5)  
grid()
```



First Model: Normal hierarchical model

The first model, suggested by WinBugs, is essentially a random effects linear growth curve

$$\begin{aligned}Y_{ij} &\sim \text{Normal}(\alpha_i + \beta_i(x_j), \tau_c) \\ \alpha_i &\sim \text{Normal}(\alpha_c, \tau_\alpha) \\ \beta_i &\sim \text{Normal}(\beta_c, \tau_\beta) \\ \alpha_c &\sim \text{Normal}(0, 1.0E - 6) \\ \beta_c &\sim \text{Normal}(0, 1.0E - 6) \\ \tau_\alpha &\sim \text{Gamma}(1.0E - 3, 1.0E - 3) \\ \tau_\beta &\sim \text{Gamma}(1.0E - 3, 1.0E - 3) \\ \tau_c &\sim \text{Gamma}(1.0E - 3, 1.0E - 3)\end{aligned}$$

where \bar{x} and τ represent the precision of Normal distribution. $\alpha_c, \beta_c, \tau_\alpha, \tau_\beta, \tau_c$ are “non-informative” priors. In the WinBugs guidelines of this dataset, the x_j is standandized around their mean in order to reduce dependence between α_i and β_i in their likelihood. We’ll not do it and we will use the priors mentioned above.

```

rats.model <- function() {

  for (i in 1:N) {

    for (j in 1:T) {

      Y[i,j] ~ dnorm(mu[i,j],tau.c)
      mu[i,j] <- alpha[i] + beta[i]*(x[j]);

    }

    alpha[i] ~ dnorm(alpha.c,tau.alpha);
    beta[i] ~ dnorm(beta.c,tau.beta); }

  alpha.c ~ dnorm(0,1.0E-6);
  beta.c ~ dnorm(0,1.0E-6);
  tau.c ~ dgamma(1.0E-3,1.0E-3);
  tau.alpha ~ dgamma(1.0E-3,1.0E-3);
  tau.beta ~ dgamma(1.0E-3,1.0E-3);
  sigma <- 1.0 / sqrt(tau.c)
  x.bar <- mean(x[]);
  alpha0 <- alpha.c - beta.c * x.bar
}

```

Now we can define the vectors of the data matrix, the starting values and the name of the parameter for JAGS:

```

rats.data <- list("Y", "x", "T", "N")
rats.params <- c("tau.c", "alpha.c", "beta.c", "tau.alpha", "tau.beta")

## Define the starting values for JAGS

rats.inits <- function(){
  list(alpha = c(250, 250, 250, 250, 250, 250, 250, 250, 250, 250, 250, 250, 250, 250, 250,
                250, 250, 250, 250, 250, 250, 250, 250, 250, 250, 250, 250, 250, 250),
        beta = c(6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6),
        alpha.c = 150, beta.c = 10,
        tau.c = 1, tau.alpha = 1, tau.beta = 1)
}

```

At this point we can run our JAGS function:

```

ratsfit <- jags(data=rats.data, inits=rats.inits, rats.params, n.chains=2, n.iter=10000, n.burnin=1000,

## module glm loaded

## Compiling model graph
##   Resolving undeclared variables
##   Allocating nodes
## Graph information:
##   Observed stochastic nodes: 150
##   Unobserved stochastic nodes: 65
##   Total graph size: 532
##
## Initializing model

```

```
ratsfit.mcmc <- as.mcmc(ratsfit)
```

```
summary(ratsfit.mcmc)
```

```
##
## Iterations = 1001:10000
## Thinning interval = 1
## Number of chains = 2
## Sample size per chain = 9000
##
## 1. Empirical mean and standard deviation for each variable,
##    plus standard error of the mean:
##
##              Mean          SD Naive SE Time-series SE
## alpha.c    106.53098  2.283557 1.702e-02    3.506e-02
## beta.c      6.18762  0.105316 7.850e-04    1.445e-03
## deviance   969.21493 14.400278 1.073e-01    2.947e-01
## tau.alpha   0.01005  0.003790 2.825e-05    7.833e-05
## tau.beta    4.28993  1.489280 1.110e-02    2.526e-02
## tau.c       0.02700  0.004012 2.990e-05    6.311e-05
##
## 2. Quantiles for each variable:
##
##              2.5%       25%       50%       75%       97.5%
## alpha.c    1.020e+02 1.050e+02 1.065e+02 108.02361 111.06482
## beta.c      5.980e+00 6.120e+00 6.189e+00  6.25767  6.39527
## deviance   9.434e+02 9.591e+02 9.683e+02 978.42834 999.37132
## tau.alpha   4.634e-03 7.374e-03 9.403e-03  0.01192  0.01923
## tau.beta    2.076e+00 3.252e+00 4.076e+00  5.07471  7.72017
## tau.c       1.967e-02 2.420e-02 2.683e-02  0.02960  0.03541
```

We can compare the results of our first model with the frequentist approach used above:


```
# INTERCEPT MEAN MODEL 1
ratsfit$BUGSoutput$mean$alpha.c
```

```
## [1] 106.531
```

```
# SLOP MEAN MODEL 2
ratsfit$BUGSoutput$mean$beta.c
```


```
## [1] 6.187619
```

```
# MODEL 1 vs. FREQ APPROACH
plot(x,colMeans(Y), lwd=4, xlab = "age (days)", ylab = "weight",
     col="red", ylim=c(135,355))
points(rep(x[1],N), Y[,1])
points(rep(x[2],N), Y[,2])
points(rep(x[3],N), Y[,3])
points(rep(x[4],N), Y[,4])
points(rep(x[5],N), Y[,5])
abline(ratsfit$BUGSoutput$mean$alpha.c,
       ratsfit$BUGSoutput$mean$beta.c, col="green", lwd=8)
abline(mean_alpha, mean_beta, col="orchid", lwd=2)
grid()
```



Another graphic method of diagnostic for convergence that we have seen at lessons, is **Gelman's plot**. This method can be applied when we have more than one chain and consists in calculating the 'within variance' of each chain, subtract from this value the 'between variance' of two chain and then multiply by a scaling factor.

```
gelman.plot(ratsfit.mcmc)
```



Second Model: Uniform priors.

In this second model we want try to change some prior parameters and see what happen at our model. Our prior sigma.alpha and sigma.beta are distributed as uniform. So we will be:

$$\alpha_i \sim Normal(\alpha_c, \tau_\alpha)$$

$$\beta_i \sim Normal(\beta_c, \tau_\beta)$$

$$\alpha_c \sim Normal(0, 1.0E - 6)$$

$$\beta_c \sim Normal(0, 1.0E - 6)$$

$$\tau_\alpha, \tau_\beta, \tau_c \sim Unif(0, 100)$$

```
rats.model2 <- function() {  
  
  for (i in 1:N) {  
  
    for (j in 1:T) {  
  
      mu[i,j] <- alpha[i] + beta[i]*(x[j]);  
      Y[i,j] ~ dnorm(mu[i,j],tau.c)  
  
    }  
  
    alpha[i] ~ dnorm(alpha.c,tau.alpha);  
    beta[i] ~ dnorm(beta.c,tau.beta); }  
  
    alpha.c ~ dnorm(0,1.0E-6);  
    beta.c ~ dnorm(0,1.0E-6);  
    tau.c <- 1.0 / (sigma*sigma);  
    sigma ~ dunif(0,100)  
    tau.alpha <- 1.0 / (sigma.alpha*sigma.alpha);  
    sigma.alpha ~ dunif(0,100)  
    tau.beta <- 1/(sigma.beta*sigma.beta)  
    sigma.beta ~ dunif(0,100)  
    x.bar <- mean(x[]);  
    alpha0 <- alpha.c - beta.c * x.bar  
  }  
  
  rats.data2 <- list("Y", "x", "T", "N")  
  rats.params2 <- c("alpha.c", "beta.c", "sigma", "sigma.alpha", "sigma.beta")  
  
  ## Define the starting values for JAGS  
  
  rats.inits2 <- function(){  
    list(alpha = c(250, 250, 250, 250, 250, 250, 250, 250, 250, 250, 250, 250, 250, 250, 250,  
                  250, 250, 250, 250, 250, 250, 250, 250, 250, 250, 250, 250, 250, 250, 250),  
         beta = c(6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6,  
                  6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6),
```



```

    alpha.c = 150, beta.c = 10,
    sigma = 1, sigma.alpha = 1, sigma.beta = 1)
}

ratsfit2 <- jags(data=rats.data2, inits=rats.inits2, rats.params2, n.chains=2, n.iter=10000, n.burnin=1000)

## Compiling model graph
##   Resolving undeclared variables
##   Allocating nodes
## Graph information:
##   Observed stochastic nodes: 150
##   Unobserved stochastic nodes: 65
##   Total graph size: 536
##
## Initializing model

ratsfit2.mcmc <- as.mcmc(ratsfit2)

summary(ratsfit2.mcmc)

##
## Iterations = 1001:10000
## Thinning interval = 1
## Number of chains = 2
## Sample size per chain = 9000
##
## 1. Empirical mean and standard deviation for each variable,
##    plus standard error of the mean:
##
##           Mean      SD Naive SE Time-series SE
## alpha.c    106.5238  2.3499 0.0175149      0.035242
## beta.c      6.1875  0.1083 0.0008073      0.001585
## deviance   969.8290 14.5171 0.1082042      0.310270
## sigma       6.1720  0.4679 0.0034878      0.008088
## sigma.alpha 10.7490  1.9566 0.0145839      0.039816
## sigma.beta  0.5153  0.0887 0.0006611      0.001587
##
## 2. Quantiles for each variable:
##
##           2.5%      25%      50%      75%      97.5%
## alpha.c    101.8581 104.9824 106.5360 108.0898 111.1189
## beta.c      5.9769  6.1149  6.1872  6.2586  6.4025
## deviance   943.6416 959.4107 969.1324 979.3029 1000.4433
## sigma       5.3411  5.8415  6.1410  6.4706  7.1690
## sigma.alpha 7.3373  9.3898 10.6034 11.9566 15.0640
## sigma.beta  0.3634  0.4537  0.5082  0.5673  0.7103

Again we can compare the results obtained by our second model with the frequentist approach computed
above:

# INTERCEPT MEAN MODEL 2
ratsfit2$BUGSoutput$mean$alpha.c

## [1] 106.5238

```


```

# SLOPE MEAN MODEL 2
ratsfit2$BUGSoutput$mean$beta.c

## [1] 6.187475

# MODEL 2 vs. FREQ APPROACH
plot(x,colMeans(Y), lwd=4, xlab = "age (days)", ylab = "weight",
     col="red", ylim=c(135,355))
points(rep(x[1],N), Y[,1])
points(rep(x[2],N), Y[,2])
points(rep(x[3],N), Y[,3])
points(rep(x[4],N), Y[,4])
points(rep(x[5],N), Y[,5])
abline(ratsfit2$BUGSoutput$mean$alpha.c,
       ratsfit2$BUGSoutput$mean$beta.c, col="grey", lwd=8)
abline(mean_alpha, mean_beta, col="orchid", lwd=2)
grid()

```



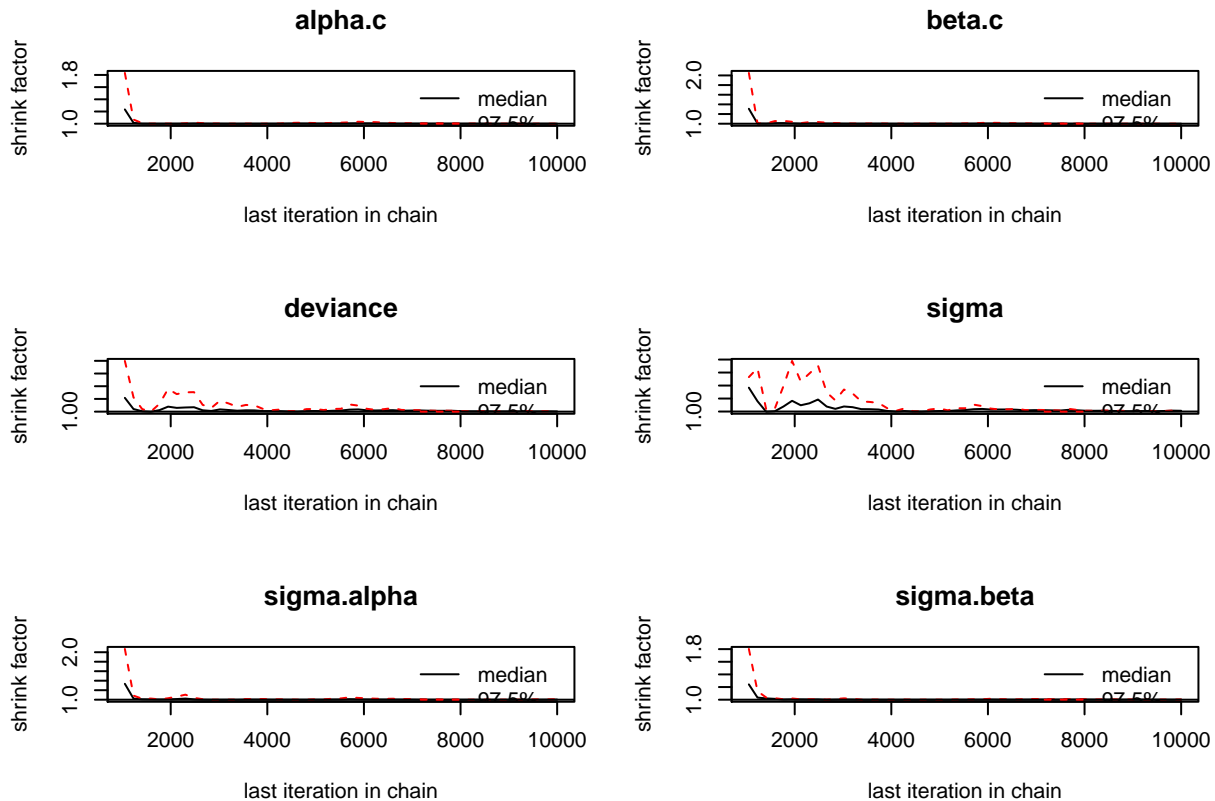
and

again we can compute the Gelman's plot:

```

gelman.plot(ratsfit2.mcmc)

```



Global comparison between models:

One way to analyze and compare our models is the **deviance information criterion** that analyze the model in terms of number of parameters and deviance. In general we prefer a model with **lower DIC**. We can consider the DIC a kind of penalized deviance, computed as follows:

$$DIC = pD + \hat{D}$$

Where:

$$pD = \text{penalty}$$

$$\hat{D} = \text{mean deviance}$$

In the previous models we have:

```
ratsfit$BUGSoutput$DIC
```

```
## [1] 1072.826
```

```
ratsfit2$BUGSoutput$DIC
```

```
## [1] 1075.149
```

The performances of each model are very influenced by number of iteration and by the burn-in parameter. We can simulate how DIC vary as the number of iteration increase with 3 different percentage of burn-in: 10%, 25% and 50%

For this section of code, output is disabled in order to avoid printing the same JAGS initialization

```
n_rep <- 100
iter <- 2000
```

```

DICmodel1_M = list()
DICmodel2_M = list()
DICmodel1 = rep(0,n_rep)
DICmodel2 = rep(0,n_rep)

for (j in 1:3){

  burn_in = c(10,25,50)
  perc = burn_in[j]/100
  DICmodel1 = rep(0,n_rep)
  DICmodel2 = rep(0,n_rep)

  for (i in 1:n_rep){

    ni=iter*i/n_rep
    nb = perc*ni

    ratsfit_SIM <- jags(data=rats.data, inits=rats.inits, rats.params, n.chains=2, n.iter=ni, n.burnin=nb)
    ratsfit2_SIM <- jags(data=rats.data2, inits=rats.inits2, rats.params2, n.chains=2, n.iter=ni, n.burnin=nb)

    DICmodel1[i] = ratsfit_SIM$BUGSoutput$DIC
    DICmodel2[i] = ratsfit2_SIM$BUGSoutput$DIC

  }

  DICmodel1_M = c(DICmodel1_M,list(DICmodel1))
  DICmodel2_M = c(DICmodel2_M,list(DICmodel2))
}

xfit = (1:n_rep)*iter/n_rep

```

Finally we can plot the results! Our DIC simulation of first model is represented by orange line and our DIC simulation of second model is represented by light blue line. Although the DIC computed in the previous step by BUGSoutput for our models is very similar (around 1070 for both models), from these simulations we can see how the value of the DIC changes in different way for our models when the number of iteration and burn-in values change.


```

yrange<-c(min(DICmodel1_M[[1]],DICmodel1_M[[2]],DICmodel1_M[[3]]),1500)

plot(xfit,DICmodel1_M[[1]],lwd=2,type="l",ylim=yrange,col='orange',ylab='DIC', xlab = 'Number of iterations',
lines(xfit,DICmodel2_M[[1]],lwd=2,type="l",col='lightseagreen')


```

10% of burn-in



```
plot(xfit,DICmodel1_M[[2]],lwd=2,type="l",ylim=yrange,col='orange',ylab='DIC', xlab = 'Number of iterations')
lines(xfit,DICmodel2_M[[2]],lwd=2,type="l",col='lightseagreen')
```

25% of burn-in



```
plot(xfit,DICmodel1_M[[3]],lwd=2,type="l",ylim=yrange,col='orange',ylab='DIC', xlab = 'Number of iterations')
lines(xfit,DICmodel2_M[[3]],lwd=2,type="l",col='lightseagreen')
```

50% of burn-in

